

Set ADT

`std::set`

`std::unordered_set`

Introduction

- In computer science, a set is an abstract data type that can store unique values, **without any particular order.**
- It is a computer implementation of the mathematical concept of a finite set.
- Unlike most other collection types, rather than retrieving a specific element from a set, one typically tests a value for membership in a set.
- Some set data structures are designed for static or frozen sets that do not change after they are constructed.
- Static sets allow only query operations on their elements — such as checking whether a given value is in the set, or enumerating the values in some arbitrary order.
- Other variants, called dynamic or mutable sets, allow also the insertion and deletion of elements from the set.

Set implementation

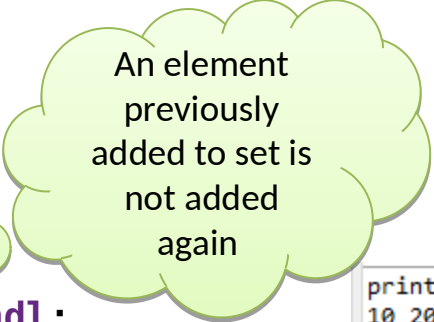
- Sets can be implemented using various data structures, which provide different time and space trade-offs for various operations.
- Some implementations are designed to improve the efficiency of very specialized operations, such as nearest or union.
- Implementations described as "general use" typically strive to optimize the **element_of**, **add**, and **delete** operations.
- A simple implementation is to use a **list**, ignoring the order of the elements and taking care to avoid repeated values. This is simple but inefficient, as operations like set membership or element deletion are $O(n)$, as they require scanning the entire list.
- Sets are often instead implemented using more efficient data structures, particularly various kinds of **trees**, or **hash tables**.

C++ `std::set` data structure in STL

- Container properties
- Sets are containers that store unique elements following a specific order.
- In a set, the value of an element also identifies it (the value is itself the key, of type T), and **each value must be unique**. The value of the elements in a set cannot be modified once in the container (the elements are always const), but they can be inserted or removed from the container.
- Internally, the elements in a set are always sorted following a specific strict weak ordering criterion indicated by its internal comparison object (of type Compare).
- **`std::set` containers are generally slower than `std::unordered_set` containers** to access individual elements by their key, but they allow the direct iteration on subsets based on their order.
- Sets are typically implemented as binary search trees.

C++ std::set data structure

```
int main() {  
    // empty set container  
    set<int> s1;  
    // insert elements in random order  
    s1.insert(40);  
    s1.insert(30);  
    s1.insert(60);  
    s1.insert(20);  
    s1.insert(50);  
    s1.insert(10);  
    cout<<"print set:"<<endl;  
    print(s1);  
  
    cout<<"\ninsert(50):"<<endl;  
    s1.insert(50);  
    print(s1);  
  
    cout<<"\nerase(30):"<<endl;  
    s1.erase(30);  
    print(s1);  
  
    cout<<"\ncopy set: "<<endl;  
    set<int> s2(s1.begin(), s1.end());  
    print(s2);  
}
```



An element
previously
added to set is
not added
again

```
void print(set<int> s1){  
    // printing set s1  
    set<int>::iterator itr;  
  
    cout << "\nThe set is : \n";  
    for (itr = s1.begin(); itr != s1.end(); itr+  
        +) {  
        cout << *itr << " ";  
    }  
    cout << endl;  
}
```

```
print set:  
10 20 30 40 50 60  
  
insert(50):  
10 20 30 40 50 60  
  
erase(30):  
10 20 40 50 60  
  
copy set:  
10 20 40 50 60
```

C++ std::set data structure

```
cout << "\nremove all elements up to 40 in s2"<<endl;
s2.erase(s2.begin(), s2.find(40));
print(s2);
cout << "\ns2.insert(90)"<<endl;
s2.insert(90);
print(s2);
cout<<"\nunion of two sets:"<<endl;
set<int> result=s1;
result.insert(s2.begin(),s2.end());
print(result);
```

```
cout<<"\n
s1.erase(s2.lower_bound(20),s1.upper_bound(50)):"<<endl;
s1.erase(s1.lower_bound(20),s1.upper_bound(50));
print(s1);
```

```
remove all elements up to 40 in s2
40 50 60
```

```
s2.insert(90)
40 50 60 90
```

```
union of two sets:
10 20 40 50 60 90
```

```
s1.erase(s2.lower_bound(20),s1.upper_bound(50)):
10 60
```

C++ std::unordered_set data structure

```
//empty unordered_set container
```

```
unordered_set<int> u1;
```

```
// insert elements in random order
```

```
u1.insert(40);
```

```
u1.insert(30);
```

```
u1.insert(60);
```

```
u1.insert(20);
```

```
u1.insert(50);
```

```
u1.insert(10);
```

```
cout<<"\nelements of
```

```
unordered_set"<<endl;
```

```
print(u1);
```

```
elements of unordered_set
```

```
10 30 40 60 20 50
```

set vs unordered_set

	set	unordered_set
Ordering	increasing order (by default)	no ordering
Implementation	Self balancing BST like <u>Red-Black Tree</u>	Hash Table
search time	$\log(n)$	$O(1)$ -> Average $O(n)$ -> Worst Case
Insertion time	$\log(n)$ + Rebalance	Same as search
Deletion time	$\log(n)$ + Rebalance	Same as search

Use set when

- We need ordered data.
- We would have to print/access the data (in sorted order).
- We need predecessor/successor of elements.
- Since set is ordered, we can use functions like `binary_search()`, `lower_bound()` and `upper_bound()` on set elements. These functions cannot be used on `unordered_set()`.

Use unordered_set when

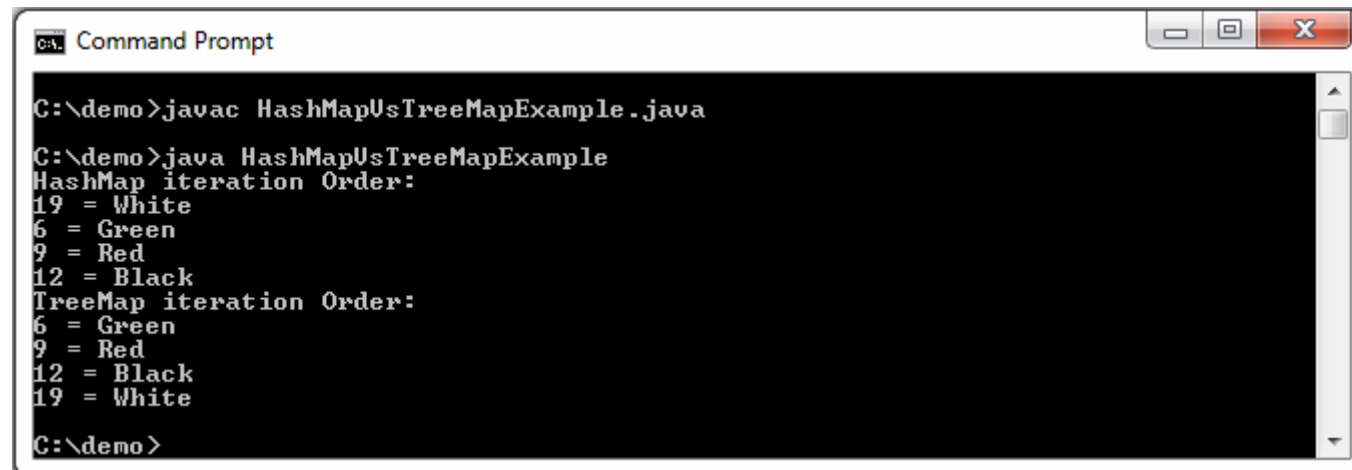
- We need to keep a set of distinct elements and no ordering is required.
- We need single element access i.e. no traversal.

map vs unordered_map

	map	unordered_map
Ordering	increasing order (by default)	no ordering
Implementation	Self balancing BST like Red-Black Tree	Hash Table
search time	$\log(n)$ 	$O(1)$ -> Average $O(n)$ -> Worst Case
Insertion time	$\log(n)$ + Rebalance	Same as search
Deletion time	$\log(n)$ + Rebalance	Same as search

Java: HashMap vs TreeMap

- The HashMap class uses the hash table.
- TreeMap internally uses a Red-Black tree, which is a self-balancing Binary Search Tree.
- HashMap: it does not maintain any order.
- TreeMap: The elements are sorted in natural order (ascending).



```
Command Prompt
C:\demo>javac HashMapVsTreeMapExample.java
C:\demo>java HashMapVsTreeMapExample
HashMap iteration Order:
19 = White
6 = Green
9 = Red
12 = Black
TreeMap iteration Order:
6 = Green
9 = Red
12 = Black
19 = White
C:\demo>
```

Java: HashSet vs TreeSet

- For operations like search, insert, and delete HashSet takes constant time for these operations on average. **HashSet is faster than TreeSet.**
- HashSet is Implemented using a hash table.
- **TreeSet takes $O(\log n)$ for search, insert and delete which is higher than HashSet.**
- **But TreeSet keeps sorted data.** Also, it supports operations like `higher()` (Returns least higher element), `floor()`, `ceiling()`, etc.
- These operations are also $O(\log n)$ in TreeSet and not supported in HashSet.
- TreeSet is implemented using a self-balancing binary search tree (Red-Black Tree). TreeSet is backed by TreeMap in Java.